

Lecture 2: Turing Machines

Our Model of Computation: Turing Machines

- In his famous 1936 paper¹, Alan Turing proposed his own **model of computation**, a.k.a. **Turing Machines** (TMs)

¹"On Computable Numbers, with an Application to the Entscheidungsproblem"

Our Model of Computation: Turing Machines

- In his famous 1936 paper¹, Alan Turing proposed his own **model of computation**, a.k.a. **Turing Machines** (TMs)
- This was an attempt to formally capture all the functions that can be computed by human “computers” following a well-defined set of rules

¹“On Computable Numbers, with an Application to the Entscheidungsproblem”

Our Model of Computation: Turing Machines

- In his famous 1936 paper¹, Alan Turing proposed his own **model of computation**, a.k.a. **Turing Machines** (TMs)
- This was an attempt to formally capture all the functions that can be computed by human “computers” following a well-defined set of rules
- Part of the impetus for the drive to define what is computable came from the mathematician David Hilbert

¹“On Computable Numbers, with an Application to the Entscheidungsproblem”

Our Model of Computation: Turing Machines

- In his famous 1936 paper¹, Alan Turing proposed his own **model of computation**, a.k.a. **Turing Machines** (TMs)
- This was an attempt to formally capture all the functions that can be computed by human “computers” following a well-defined set of rules
- Part of the impetus for the drive to define what is computable came from the mathematician David Hilbert
 - Hilbert wondered if it exists an “effective procedure” (i.e., our informal definition of algorithm) that decides whether any mathematical statement is true or false, in a finite number of steps

¹“On Computable Numbers, with an Application to the Entscheidungsproblem”

Our Model of Computation: Turing Machines

- In his famous 1936 paper¹, Alan Turing proposed his own **model of computation**, a.k.a. **Turing Machines** (TMs)
- This was an attempt to formally capture all the functions that can be computed by human “computers” following a well-defined set of rules
- Part of the impetus for the drive to define what is computable came from the mathematician David Hilbert
 - Hilbert wondered if it exists an “effective procedure” (i.e., our informal definition of algorithm) that decides whether any mathematical statement is true or false, in a finite number of steps
 - As a special case of this decision problem, Hilbert considered the validity problem for first-order logic (a.k.a. *entscheidungsproblem*)

¹“On Computable Numbers, with an Application to the Entscheidungsproblem”

Turing Machines: An Informal Perspective

- To describe his machine, Turing thought of a person as having access to as much “paper” as they need (i.e., simulating infinite memory)

Turing Machines: An Informal Perspective

- To describe his machine, Turing thought of a person as having access to as much “paper” as they need (i.e., simulating infinite memory)
- We can think of this paper as a one-dimensional **tape**

Turing Machines: An Informal Perspective

- To describe his machine, Turing thought of a person as having access to as much “paper” as they need (i.e., simulating infinite memory)
- We can think of this paper as a one-dimensional **tape**
- The tape is divided into “cells”, where each cell can hold a single symbol (i.e., some element of a finite alphabet)

Turing Machines: An Informal Perspective

- To describe his machine, Turing thought of a person as having access to as much “paper” as they need (i.e., simulating infinite memory)
- We can think of this paper as a one-dimensional **tape**
- The tape is divided into “cells”, where each cell can hold a single symbol (i.e., some element of a finite alphabet)
- At any point in time, the person can read from and write to a single cell of the paper

Turing Machines: An Informal Perspective

- To describe his machine, Turing thought of a person as having access to as much “paper” as they need (i.e., simulating infinite memory)
- We can think of this paper as a one-dimensional **tape**
- The tape is divided into “cells”, where each cell can hold a single symbol (i.e., some element of a finite alphabet)
- At any point in time, the person can read from and write to a single cell of the paper
- Based on the content of a cell, the person can update their finite (mental) state, and/or move to the cell immediately to the left or right of the current one

Turing Machines: An Informal Perspective

- To describe his machine, Turing thought of a person as having access to as much “paper” as they need (i.e., simulating infinite memory)
- We can think of this paper as a one-dimensional **tape**
- The tape is divided into “cells”, where each cell can hold a single symbol (i.e., some element of a finite alphabet)
- At any point in time, the person can read from and write to a single cell of the paper
- Based on the content of a cell, the person can update their finite (mental) state, and/or move to the cell immediately to the left or right of the current one

Note

The linear nature of memory tape, as opposed to random access memory, is a limitation on computation speed but not power: a TM can find any memory location, i.e., tape cell, by sequentially scanning its tape

Turing Machines: A Formal Definition

Definition (Turing machine)

A Turing machine M is a 6-tuple $(Q, \Sigma, \delta_M, q_0, q_{\text{accept}}, q_{\text{reject}})$, where:

- Q is the finite set of **states**

Turing Machines: A Formal Definition

Definition (Turing machine)

A Turing machine M is a 6-tuple $(Q, \Sigma, \delta_M, q_0, q_{\text{accept}}, q_{\text{reject}})$, where:

- Q is the finite set of **states**
- $\Sigma \cup \{\emptyset\}$, where Σ is the finite **alphabet** and \emptyset is the special **blank symbol**

Turing Machines: A Formal Definition

Definition (Turing machine)

A Turing machine M is a 6-tuple $(Q, \Sigma, \delta_M, q_0, q_{\text{accept}}, q_{\text{reject}})$, where:

- Q is the finite set of **states**
- $\Sigma \cup \{\emptyset\}$, where Σ is the finite **alphabet** and \emptyset is the special **blank symbol**
- $\delta_M : Q \times \Sigma \cup \{\emptyset\} \mapsto Q \times \Sigma \cup \{\emptyset\} \times \{-1, 0, +1\}$ is the **transition function**

Turing Machines: A Formal Definition

Definition (Turing machine)

A Turing machine M is a 6-tuple $(Q, \Sigma, \delta_M, q_0, q_{\text{accept}}, q_{\text{reject}})$, where:

- Q is the finite set of **states**
- $\Sigma \cup \{\emptyset\}$, where Σ is the finite **alphabet** and \emptyset is the special **blank symbol**
- $\delta_M : Q \times \Sigma \cup \{\emptyset\} \mapsto Q \times \Sigma \cup \{\emptyset\} \times \{-1, 0, +1\}$ is the **transition function**
- $q_0 \in Q$ is the **start state**

Turing Machines: A Formal Definition

Definition (Turing machine)

A Turing machine M is a 6-tuple $(Q, \Sigma, \delta_M, q_0, q_{\text{accept}}, q_{\text{reject}})$, where:

- Q is the finite set of **states**
- $\Sigma \cup \{\emptyset\}$, where Σ is the finite **alphabet** and \emptyset is the special **blank symbol**
- $\delta_M : Q \times \Sigma \cup \{\emptyset\} \mapsto Q \times \Sigma \cup \{\emptyset\} \times \{-1, 0, +1\}$ is the **transition function**
- $q_0 \in Q$ is the **start state**
- $q_{\text{accept}} \in Q$ is the **accept state**

Turing Machines: A Formal Definition

Definition (Turing machine)

A Turing machine M is a 6-tuple $(Q, \Sigma, \delta_M, q_0, q_{\text{accept}}, q_{\text{reject}})$, where:

- Q is the finite set of **states**
- $\Sigma \cup \{\emptyset\}$, where Σ is the finite **alphabet** and \emptyset is the special **blank symbol**
- $\delta_M : Q \times \Sigma \cup \{\emptyset\} \mapsto Q \times \Sigma \cup \{\emptyset\} \times \{-1, 0, +1\}$ is the **transition function**
- $q_0 \in Q$ is the **start state**
- $q_{\text{accept}} \in Q$ is the **accept state**
- $q_{\text{reject}} \in Q$ is the **reject state**, s.t. $q_{\text{accept}} \neq q_{\text{reject}}$

Turing Machine: How Does It Work?

- The generic Turing machine M receives its input on the tape, i.e., $\sigma_{\text{in}} \in \Sigma^*$

Turing Machine: How Does It Work?

- The generic Turing machine M receives its input on the tape, i.e., $\sigma_{\text{in}} \in \Sigma^*$
- The rest of the tape is filled with blank symbols (\emptyset)

Turing Machine: How Does It Work?

- The generic Turing machine M receives its input on the tape, i.e., $\sigma_{\text{in}} \in \Sigma^*$
- The rest of the tape is filled with blank symbols (\emptyset)
- The read/write head starts from the leftmost cell of the tape (e.g., by convention indicated with position $h = 0$)

Turing Machine: How Does It Work?

- The generic Turing machine M receives its input on the tape, i.e., $\sigma_{\text{in}} \in \Sigma^*$
- The rest of the tape is filled with blank symbols (\emptyset)
- The read/write head starts from the leftmost cell of the tape (e.g., by convention indicated with position $h = 0$)
- The head keeps scanning the cells on the tape, according to the transition function δ_M

Turing Machine: How Does It Work?

- The generic Turing machine M receives its input on the tape, i.e., $\sigma_{\text{in}} \in \Sigma^*$
- The rest of the tape is filled with blank symbols (\emptyset)
- The read/write head starts from the leftmost cell of the tape (e.g., by convention indicated with position $h = 0$)
- The head keeps scanning the cells on the tape, according to the transition function δ_M
- The computation continues until it either enters q_{accept} or q_{reject} state, otherwise M may run forever

Turing Machine: How Does It Work?

- The generic Turing machine M receives its input on the tape, i.e., $\sigma_{\text{in}} \in \Sigma^*$
- The rest of the tape is filled with blank symbols (\emptyset)
- The read/write head starts from the leftmost cell of the tape (e.g., by convention indicated with position $h = 0$)
- The head keeps scanning the cells on the tape, according to the transition function δ_M
- The computation continues until it either enters q_{accept} or q_{reject} state, otherwise M may run forever
- If M ever halts, it will leave the output string on the tape, i.e., $\sigma_{\text{out}} \in \Sigma^*$

Turing Machine: The Transition Function δ_M

- From any given state $q \in Q \setminus \{q_{\text{accept}}, q_{\text{reject}}\}$ and the symbol in the current head position h , δ_M specifies:

Turing Machine: The Transition Function δ_M

- From any given state $q \in Q \setminus \{q_{\text{accept}}, q_{\text{reject}}\}$ and the symbol in the current head position h , δ_M specifies:
 - the new state $q' \in Q$ that M should enter

Turing Machine: The Transition Function δ_M

- From any given state $q \in Q \setminus \{q_{\text{accept}}, q_{\text{reject}}\}$ and the symbol in the current head position h , δ_M specifies:
 - the new state $q' \in Q$ that M should enter
 - the new symbol $\sigma' \in \Sigma \cup \{\emptyset\}$ that should be written in the current tape cell

Turing Machine: The Transition Function δ_M

- From any given state $q \in Q \setminus \{q_{\text{accept}}, q_{\text{reject}}\}$ and the symbol in the current head position h , δ_M specifies:
 - the new state $q' \in Q$ that M should enter
 - the new symbol $\sigma' \in \Sigma \cup \{\emptyset\}$ that should be written in the current tape cell
 - the new head position, $h' = h + d$, where $d \in \{-1, 0, 1\}$

Turing Machine: The Transition Function δ_M

- From any given state $q \in Q \setminus \{q_{\text{accept}}, q_{\text{reject}}\}$ and the symbol in the current head position h , δ_M specifies:
 - the new state $q' \in Q$ that M should enter
 - the new symbol $\sigma' \in \Sigma \cup \{\emptyset\}$ that should be written in the current tape cell
 - the new head position, $h' = h + d$, where $d \in \{-1, 0, 1\}$

Note

One should not confuse the transition function δ_M of a Turing machine M with the function f_M that the machine computes:

Turing Machine: The Transition Function δ_M

- From any given state $q \in Q \setminus \{q_{\text{accept}}, q_{\text{reject}}\}$ and the symbol in the current head position h , δ_M specifies:
 - the new state $q' \in Q$ that M should enter
 - the new symbol $\sigma' \in \Sigma \cup \{\emptyset\}$ that should be written in the current tape cell
 - the new head position, $h' = h + d$, where $d \in \{-1, 0, 1\}$

Note

One should not confuse the transition function δ_M of a Turing machine M with the function f_M that the machine computes:

- δ_M is a **finite** function, which takes $|Q| \cdot |\Sigma \cup \{\emptyset\}|$ possible inputs and produces $3 \cdot |Q| \cdot |\Sigma \cup \{\emptyset\}|$ possible outputs

Turing Machine: The Transition Function δ_M

- From any given state $q \in Q \setminus \{q_{\text{accept}}, q_{\text{reject}}\}$ and the symbol in the current head position h , δ_M specifies:
 - the new state $q' \in Q$ that M should enter
 - the new symbol $\sigma' \in \Sigma \cup \{\emptyset\}$ that should be written in the current tape cell
 - the new head position, $h' = h + d$, where $d \in \{-1, 0, 1\}$

Note

One should not confuse the transition function δ_M of a Turing machine M with the function f_M that the machine computes:

- δ_M is a **finite** function, which takes $|Q| \cdot |\Sigma \cup \{\emptyset\}|$ possible inputs and produces $3 \cdot |Q| \cdot |\Sigma \cup \{\emptyset\}|$ possible outputs
- The machine can compute an **infinite** function f_M that takes as input a string $\sigma_{\text{in}} \in \Sigma^*$ and produces another string $\sigma_{\text{out}} \in \Sigma^*$ as output, both of arbitrary lengths

From Turing Machines to Computable Functions

Definition (Computable Function)

Let $f : \Sigma^* \mapsto \Sigma^*$ be a (total) function and let M be a Turing machine. We say that M computes f if for every $x \in \Sigma^*$, $M(x) = f(x)$.

We say that a function f is computable if there exists a Turing machine M that computes it.

From Turing Machines to Computable Functions

Definition (Computable Function)

Let $f : \Sigma^* \mapsto \Sigma^*$ be a (total) function and let M be a Turing machine. We say that M computes f if for every $x \in \Sigma^*$, $M(x) = f(x)$.

We say that a function f is computable if there exists a Turing machine M that computes it.

Note

Defining a function “computable” if it can be computed by a Turing machine might seem incautious, but this is equivalent to being computable in virtually *any* reasonable model of computation.

The Church-Turing Thesis

- A hypothesis about the nature of computable functions

The Church-Turing Thesis

- A hypothesis about the nature of computable functions
- A function can be calculated by an effective method if and only if it is computable by a Turing machine
 - Or by any equivalent computational models proposed by Gödel (**recursive functions**) and Church (**λ -calculus**)

The Church-Turing Thesis

- A hypothesis about the nature of computable functions
- A function can be calculated by an effective method if and only if it is computable by a Turing machine
 - Or by any equivalent computational models proposed by Gödel (**recursive functions**) and Church (**λ -calculus**)
- The three formally-defined classes of computable functions coincide with the informal notion of an effectively calculable function

The Church-Turing Thesis

- A hypothesis about the nature of computable functions
- A function can be calculated by an effective method if and only if it is computable by a Turing machine
 - Or by any equivalent computational models proposed by Gödel (**recursive functions**) and Church (**λ -calculus**)
- The three formally-defined classes of computable functions coincide with the informal notion of an effectively calculable function
- Since the concept of effective calculability does not have a formal definition, the thesis, although it has near-universal acceptance, cannot be formally proven

(Boolean) Computable Functions

- Often, we are interested in functions $f : \Sigma^* \mapsto \Sigma$, i.e., those having a single bit of output

(Boolean) Computable Functions

- Often, we are interested in functions $f : \Sigma^* \mapsto \Sigma$, i.e., those having a single bit of output
- We can give a special name for the set of such boolean computable functions

(Boolean) Computable Functions

- Often, we are interested in functions $f : \Sigma^* \mapsto \Sigma$, i.e., those having a single bit of output
- We can give a special name for the set of such boolean computable functions

Definition

We define by \mathcal{R} the set of **all** computable functions $f : \Sigma^* \mapsto \Sigma$

Functions vs. Languages

- Many texts use the terminology of “languages” rather than functions to refer to computational tasks

Functions vs. Languages

- Many texts use the terminology of “languages” rather than functions to refer to computational tasks
- A Turing machine M **decides** a language L if for **every** input $x \in \Sigma^*$, $M(x)$ outputs 1 if and only if $x \in L$ (0, otherwise)

Functions vs. Languages

- Many texts use the terminology of “languages” rather than functions to refer to computational tasks
- A Turing machine M **decides** a language L if for **every** input $x \in \Sigma^*$, $M(x)$ outputs 1 if and only if $x \in L$ (0, otherwise)
- This is equivalent to computing the boolean (total) function $f : \Sigma^* \mapsto \Sigma$ defined as:

$$f(x) = \begin{cases} 1, & \text{if } x \in L \\ 0, & \text{otherwise} \end{cases}$$

Functions vs. Languages

- Many texts use the terminology of “languages” rather than functions to refer to computational tasks
- A Turing machine M **decides** a language L if for **every** input $x \in \Sigma^*$, $M(x)$ outputs 1 if and only if $x \in L$ (0, otherwise)
- This is equivalent to computing the boolean (total) function $f : \Sigma^* \mapsto \Sigma$ defined as:

$$f(x) = \begin{cases} 1, & \text{if } x \in L \\ 0, & \text{otherwise} \end{cases}$$

Definition (Turing-decidable Language)

A language L is **Turing-decidable** (or simply **decidable**) if there is a Turing machine M that decides it

A Note on the Terminology

- For historical reasons, some texts also refer to computable boolean functions/decidable languages as **recursive languages**

A Note on the Terminology

- For historical reasons, some texts also refer to computable boolean functions/decidable languages as **recursive languages**
- This is also the reason why the letter \mathcal{R} is often used

A Note on the Terminology

- For historical reasons, some texts also refer to computable boolean functions/decidable languages as **recursive languages**
- This is also the reason why the letter \mathcal{R} is often used
- We stick to the term *functions* rather than *languages*, although the following always holds:

$$f : \Sigma^* \mapsto \Sigma$$

$$L = \{x \in \Sigma^* \mid f(x) = 1\}$$

Infinite Loops and Partial Functions

- Given a Turing machine M , we cannot determine a priori the length of its output

Infinite Loops and Partial Functions

- Given a Turing machine M , we cannot determine a priori the length of its output
- In fact, we don't even know if an output will be produced at all!

Infinite Loops and Partial Functions

- Given a Turing machine M , we cannot determine a priori the length of its output
- In fact, we don't even know if an output will be produced at all!
 - For example, it is easy to design a Turing machine whose transition function never leads to a halting state (i.e., either q_{accept} or q_{reject})

Infinite Loops and Partial Functions

- Given a Turing machine M , we cannot determine a priori the length of its output
- In fact, we don't even know if an output will be produced at all!
 - For example, it is easy to design a Turing machine whose transition function never leads to a halting state (i.e., either q_{accept} or q_{reject})
- If a machine M fails to stop and produce an output on some input $x \in \Sigma^*$, then it cannot compute any total function f

Infinite Loops and Partial Functions

- Given a Turing machine M , we cannot determine a priori the length of its output
- In fact, we don't even know if an output will be produced at all!
 - For example, it is easy to design a Turing machine whose transition function never leads to a halting state (i.e., either q_{accept} or q_{reject})
- If a machine M fails to stop and produce an output on some input $x \in \Sigma^*$, then it cannot compute any total function f
- However, M can still compute a **partial function**

Infinite Loops and Partial Functions

- Given a Turing machine M , we cannot determine a priori the length of its output
- In fact, we don't even know if an output will be produced at all!
 - For example, it is easy to design a Turing machine whose transition function never leads to a halting state (i.e., either q_{accept} or q_{reject})
- If a machine M fails to stop and produce an output on some input $x \in \Sigma^*$, then it cannot compute any total function f
- However, M can still compute a **partial function**

Definition

A partial function $f : A \mapsto B$ is a function that is only defined on a subset A' of A (i.e., $A' \subset A$). We can also think of such a function as mapping from A to $B \cup \{\perp\}$, where \perp is a special “failure” symbol such that $f(a) = \perp$ indicates f is not defined on input a

Turing Machines Computing Partial Functions

Example

Consider the function $div : \mathbb{Z}^{0+} \times \mathbb{Z}^{0+} \mapsto \mathbb{Z}^{0+}$, defined as follows:

$$div(a, b) = \begin{cases} \lceil \frac{a}{b} \rceil, & \text{if } b > 0 \\ \perp, & \text{otherwise} \end{cases}$$

Turing Machines Computing Partial Functions

Example

Consider the function $div : \mathbb{Z}^{0+} \times \mathbb{Z}^{0+} \mapsto \mathbb{Z}^{0+}$, defined as follows:

$$div(a, b) = \begin{cases} \lceil \frac{a}{b} \rceil, & \text{if } b > 0 \\ \perp, & \text{otherwise} \end{cases}$$

- We can design a Turing machine M that computes div on inputs a, b by outputting the first $c \in \{0, 1, 2, \dots\}$ such that $cb \geq a$

Turing Machines Computing Partial Functions

Example

Consider the function $div : \mathbb{Z}^{0+} \times \mathbb{Z}^{0+} \mapsto \mathbb{Z}^{0+}$, defined as follows:

$$div(a, b) = \begin{cases} \lceil \frac{a}{b} \rceil, & \text{if } b > 0 \\ \perp, & \text{otherwise} \end{cases}$$

- We can design a Turing machine M that computes div on inputs a, b by outputting the first $c \in \{0, 1, 2, \dots\}$ such that $cb \geq a$
 - If $a > 0$ and $b = 0$, M never halts but this is ok, since div is undefined on such inputs

Turing Machines Computing Partial Functions

Example

Consider the function $div : \mathbb{Z}^{0+} \times \mathbb{Z}^{0+} \mapsto \mathbb{Z}^{0+}$, defined as follows:

$$div(a, b) = \begin{cases} \lceil \frac{a}{b} \rceil, & \text{if } b > 0 \\ \perp, & \text{otherwise} \end{cases}$$

- We can design a Turing machine M that computes div on inputs a, b by outputting the first $c \in \{0, 1, 2, \dots\}$ such that $cb \geq a$
 - If $a > 0$ and $b = 0$, M never halts but this is ok, since div is undefined on such inputs
 - If $a = b = 0$, M will output 0, which is also ok, since we do not care about what the machine outputs on inputs on which div is undefined

Computable Functions (Redefined)

Definition

Let f be a **total** or **partial** function, such that $f : \Sigma^* \mapsto \Sigma^*$ and let M be a Turing machine.

We say that M **computes** f if for every $x \in \Sigma^*$ on which f is defined, $M(x) = f(x)$.

We say that a (partial or total) function f is **computable** if there is a Turing machine that computes it.

A Clarification on the Role of \perp

- We used \perp as our special “failure symbol”; if a Turing machine M fails to halt on some input $x \in \Sigma^*$ then we denote this by $M(x) = \perp$

A Clarification on the Role of \perp

- We used \perp as our special “failure symbol”; if a Turing machine M fails to halt on some input $x \in \Sigma^*$ then we denote this by $M(x) = \perp$
- This **does not** mean that M outputs some encoding of the symbol \perp but rather that M enters into an infinite loop when given x as input

A Clarification on the Role of \perp

- We used \perp as our special “failure symbol”; if a Turing machine M fails to halt on some input $x \in \Sigma^*$ then we denote this by $M(x) = \perp$
- This **does not** mean that M outputs some encoding of the symbol \perp but rather that M enters into an infinite loop when given x as input
- As such, one might be tempted to think that M halts on x if and only if f is defined on x

A Clarification on the Role of \perp

- We used \perp as our special “failure symbol”; if a Turing machine M fails to halt on some input $x \in \Sigma^*$ then we denote this by $M(x) = \perp$
- This **does not** mean that M outputs some encoding of the symbol \perp but rather that M enters into an infinite loop when given x as input
- As such, one might be tempted to think that M halts on x if and only if f is defined on x
- However, for a Turing machine M to compute a partial function f it is **not** necessary to enter an infinite loop on inputs x outside the domain of f

A Clarification on the Role of \perp

- We used \perp as our special “failure symbol”; if a Turing machine M fails to halt on some input $x \in \Sigma^*$ then we denote this by $M(x) = \perp$
- This **does not** mean that M outputs some encoding of the symbol \perp but rather that M enters into an infinite loop when given x as input
- As such, one might be tempted to think that M halts on x if and only if f is defined on x
- However, for a Turing machine M to compute a partial function f it is **not** necessary to enter an infinite loop on inputs x outside the domain of f
- All that is needed is for M to output $f(x)$ on $x \in \text{domain}(f)$: on any other input it is OK for M to output an arbitrary value or not to halt at all

Functions vs. Languages

- A Turing machine M **recognizes** a language L if for **every** input $x \in \Sigma^*$, $M(x)$ outputs 1 if and only if $x \in L$

Functions vs. Languages

- A Turing machine M **recognizes** a language L if for **every** input $x \in \Sigma^*$, $M(x)$ outputs 1 if and only if $x \in L$
- If $x \notin L$, M may either halt with non-sense output or loop forever

Functions vs. Languages

- A Turing machine M **recognizes** a language L if for **every** input $x \in \Sigma^*$, $M(x)$ outputs 1 if and only if $x \in L$
- If $x \notin L$, M may either halt with non-sense output or loop forever
- This is equivalent to computing the partial function $f : \Sigma^* \mapsto \Sigma$ defined as:

$$f(x) = \begin{cases} 1, & \text{if } x \in L \\ \perp, & \text{otherwise} \end{cases}$$

Functions vs. Languages

- A Turing machine M **recognizes** a language L if for **every** input $x \in \Sigma^*$, $M(x)$ outputs 1 if and only if $x \in L$
- If $x \notin L$, M may either halt with non-sense output or loop forever
- This is equivalent to computing the partial function $f : \Sigma^* \mapsto \Sigma$ defined as:

$$f(x) = \begin{cases} 1, & \text{if } x \in L \\ \perp, & \text{otherwise} \end{cases}$$

Definition (Turing-recognizable Language)

A language L is **Turing-recognizable** (or simply **recognizable** or **semi-decidable**) if there is a Turing machine M that recognizes it

A Note on the Terminology

- For historical reasons, some texts also refer to recognizable languages as **recursively enumerable languages**

A Note on the Terminology

- For historical reasons, some texts also refer to recognizable languages as **recursively enumerable languages**
- This is also the reason why the letter \mathcal{RE} is often used

A Note on the Terminology

- For historical reasons, some texts also refer to recognizable languages as **recursively enumerable languages**
- This is also the reason why the letter \mathcal{RE} is often used
- We stick to the term *functions* rather than *languages*, although the following always holds:

$$f : \Sigma^* \mapsto \Sigma$$

$$L = \{x \in \Sigma^* \mid f(x) = 1\}$$

Variants of Turing Machines

- Alternative definitions of Turing machines abound, e.g., multiple tapes or non-deterministic Turing machines

Variants of Turing Machines

- Alternative definitions of Turing machines abound, e.g., multiple tapes or non-deterministic Turing machines
- Interestingly enough, the original computational model and its variants have all the same power

Variants of Turing Machines

- Alternative definitions of Turing machines abound, e.g., multiple tapes or non-deterministic Turing machines
- Interestingly enough, the original computational model and its variants have all the same power
- They all compute the same functions/recognize the same set of languages

Multi-tape Turing Machines

- Like an ordinary Turing machine, yet with several tapes

Multi-tape Turing Machines

- Like an ordinary Turing machine, yet with several tapes
- Each tape has its own reading/writing head

Multi-tape Turing Machines

- Like an ordinary Turing machine, yet with several tapes
- Each tape has its own reading/writing head
- Initially, the input is located on the first tape (i.e., tape 1), whilst the others are filled with blank symbols

Multi-tape Turing Machines

- Like an ordinary Turing machine, yet with several tapes
- Each tape has its own reading/writing head
- Initially, the input is located on the first tape (i.e., tape 1), whilst the others are filled with blank symbols
- The transition function δ_M is changed to allow for reading, writing, and moving the heads on some or all of the tapes, simultaneously

Multi-tape Turing Machines

- Like an ordinary Turing machine, yet with several tapes
- Each tape has its own reading/writing head
- Initially, the input is located on the first tape (i.e., tape 1), whilst the others are filled with blank symbols
- The transition function δ_M is changed to allow for reading, writing, and moving the heads on some or all of the tapes, simultaneously
- Formally, the transition function of a k -tape Turing machine is defined as follows:

$$\delta_M : Q \times \Sigma^k \cup \{\emptyset\}^k \mapsto Q \times \Sigma^k \cup \{\emptyset\}^k \times \{-1, 0, +1\}^k$$

Multi-tape Turing Machines: Example

Consider a k -tape Turing Machine, then the expression

$$\delta_M(q_i, \sigma_1, \sigma_2, \dots, \sigma_k) = (q_j, \sigma'_1, \sigma'_2, \dots, \sigma'_k, +1, 0, \dots, -1)$$

means that, if the machine is in state q_i and heads 1 through k are reading symbols σ_1 through σ_k , then it goes to state q_j , writes symbols σ'_1 through σ'_k and moves each head to the left (-1) or to the right (+1) of the current position, or leaves it where it is (0)

Equivalence Between Single- and Multi-Tape TMs

- Intuitively, multi-tape Turing machines seem more powerful than ordinary, single-tape Turing machines

Equivalence Between Single- and Multi-Tape TMs

- Intuitively, multi-tape Turing machines seem more powerful than ordinary, single-tape Turing machines
- In fact, it can be proven that those two models of computations are indeed equivalent (i.e., they both recognize the same languages)

Equivalence Between Single- and Multi-Tape TMs

- Intuitively, multi-tape Turing machines seem more powerful than ordinary, single-tape Turing machines
- In fact, it can be proven that those two models of computations are indeed equivalent (i.e., they both recognize the same languages)
- To sketch the idea of the proof, consider two Turing machines: S , M
 - The former is a single-tape machine, whilst the latter is multi-tape
 - The key idea is to simulate M using S
 - We can lay down the content of the k tapes of M on the single tape of S , using a special symbol as delimiter (e.g., $\#$)
 - Add another extra symbol (e.g., \bullet) on top of the current symbol to mimic the head position on each tape

Non-deterministic Turing Machines (NTMs)

- At any time during the computation a non-deterministic TM proceeds according to several possibilities

Non-deterministic Turing Machines (NTMs)

- At any time during the computation a non-deterministic TM proceeds according to several possibilities
- The transition function for a NTM δ_M is defined as follows:

$$\delta_M : Q \times \Sigma \cup \{\emptyset\} \mapsto \mathcal{P}(Q \times \Sigma \cup \{\emptyset\} \times \{-1, 0, +1\})$$

where $\mathcal{P}(A)$ stands for the **power set** of A , i.e., the set of all subsets of A

Non-deterministic Turing Machines (NTMs)

- At any time during the computation a non-deterministic TM proceeds according to several possibilities
- The transition function for a NTM δ_M is defined as follows:

$$\delta_M : Q \times \Sigma \cup \{\emptyset\} \mapsto \mathcal{P}(Q \times \Sigma \cup \{\emptyset\} \times \{-1, 0, +1\})$$

where $\mathcal{P}(A)$ stands for the **power set** of A , i.e., the set of all subsets of A

- The computation of an NTM is a **tree**, whose branches correspond to different computational paths for the machine

Non-deterministic Turing Machines (NTMs)

- At any time during the computation a non-deterministic TM proceeds according to several possibilities
- The transition function for a NTM δ_M is defined as follows:

$$\delta_M : Q \times \Sigma \cup \{\emptyset\} \mapsto \mathcal{P}(Q \times \Sigma \cup \{\emptyset\} \times \{-1, 0, +1\})$$

where $\mathcal{P}(A)$ stands for the **power set** of A , i.e., the set of all subsets of A

- The computation of an NTM is a **tree**, whose branches correspond to different computational paths for the machine
- If some branch leads to the accept state (q_{accept}), the machine accepts its input

Equivalence Between Deterministic and Non-Deterministic TMs

- Again, intuitively NTMs seem more powerful than ordinary, deterministic TMs

Equivalence Between Deterministic and Non-Deterministic TMs

- Again, intuitively NTMs seem more powerful than ordinary, deterministic TMs
- In fact, it can be proven that those two models of computations are indeed equivalent (i.e., they both recognize the same languages)

Equivalence Between Deterministic and Non-Deterministic TMs

- Again, intuitively NTMs seem more powerful than ordinary, deterministic TMs
- In fact, it can be proven that those two models of computations are indeed equivalent (i.e., they both recognize the same languages)
- To sketch the idea of the proof, consider two Turing machines: D , N
 - The former is a deterministic machine, whilst the latter is non-deterministic
 - The key idea is to simulate N using D by letting D try **all** the possible branches of N 's non-deterministic computation
 - If D ever reaches the accept state on one of these branches, D accepts; otherwise D 's simulation may run forever

Non-Deterministic Computation as a Tree

- We can view N 's computation on an input string x as a tree

Non-Deterministic Computation as a Tree

- We can view N 's computation on an input string x as a tree
- Each node of such a tree is a configuration of N , with the **root node** being the starting configuration

Non-Deterministic Computation as a Tree

- We can view N 's computation on an input string x as a tree
- Each node of such a tree is a configuration of N , with the **root node** being the starting configuration
- The machine D searches this tree for an accepting configuration (i.e., a configuration whose state is q_{accept})

Non-Deterministic Computation as a Tree

- We can view N 's computation on an input string x as a tree
- Each node of such a tree is a configuration of N , with the **root node** being the starting configuration
- The machine D searches this tree for an accepting configuration (i.e., a configuration whose state is q_{accept})
- **breadth first search** explores all branches at the same depth of the tree before moving to the next level

Non-Deterministic Computation as a Tree

- We can view N 's computation on an input string x as a tree
- Each node of such a tree is a configuration of N , with the **root node** being the starting configuration
- The machine D searches this tree for an accepting configuration (i.e., a configuration whose state is q_{accept})
- **breadth first search** explores all branches at the same depth of the tree before moving to the next level
- This guarantees that D will visit every node in the tree until it encounters an accepting configuration

Universal Turing Machine

- So far, we have roughly assumed that a Turing Machine M takes as input some x encoded as a binary string and computes a function $f(x)$

Universal Turing Machine

- So far, we have roughly assumed that a Turing Machine M takes as input some x encoded as a binary string and computes a function $f(x)$
- We have already seen that we can use the same binary string encoding to represent virtually **any** object

Universal Turing Machine

- So far, we have roughly assumed that a Turing Machine M takes as input some x encoded as a binary string and computes a function $f(x)$
- We have already seen that we can use the same binary string encoding to represent virtually **any** object
- As a special case, we can therefore encode **any** Turing machine M together with **any** of its input x

Universal Turing Machine

Definition (Universal Turing Machine)

There exists a Turing machine U , such that on every string M which encodes a Turing machine, and $x \in \Sigma^*$:

$$U(M, x) = M(x)$$

If the machine M halts on x and outputs some $y \in \Sigma^*$ (i.e., $M(x) = y$), then:

$$U(M, x) = M(x) = y$$

If M does **not** halt on x (i.e., $M(x) = \perp$) then:

$$U(M, x) = M(x) = \perp$$

Universal Turing Machine: Intuition

- Intuitively, the existence of U implies the existence of a “universal” algorithm that can evaluate arbitrary algorithms (M) on arbitrary inputs (x)

Universal Turing Machine: Intuition

- Intuitively, the existence of U implies the existence of a “universal” algorithm that can evaluate arbitrary algorithms (M) on arbitrary inputs (x)
- The desired program U is an **interpreter** for Turing machines

Universal Turing Machine: Intuition

- Intuitively, the existence of U implies the existence of a “universal” algorithm that can evaluate arbitrary algorithms (M) on arbitrary inputs (x)
- The desired program U is an **interpreter** for Turing machines
- U gets a representation of the machine M (e.g., source code), and some input x , and simulates the execution of M on x

The Existence of a Universal Turing Machine

- How would you code U in your favorite programming language?

The Existence of a Universal Turing Machine

- How would you code U in your favorite programming language?
- First, you need to decide on some representation scheme for M .
 - For example, you can use an array or a dictionary to encode M 's transition function

The Existence of a Universal Turing Machine

- How would you code U in your favorite programming language?
- First, you need to decide on some representation scheme for M .
 - For example, you can use an array or a dictionary to encode M 's transition function
- Then you would use some data structure, such as a list, to store the contents of M 's tape

The Existence of a Universal Turing Machine

- How would you code U in your favorite programming language?
- First, you need to decide on some representation scheme for M .
 - For example, you can use an array or a dictionary to encode M 's transition function
- Then you would use some data structure, such as a list, to store the contents of M 's tape
- Now you can simulate M step by step, updating the data structure as you move along

The Existence of a Universal Turing Machine

- How would you code U in your favorite programming language?
- First, you need to decide on some representation scheme for M .
 - For example, you can use an array or a dictionary to encode M 's transition function
- Then you would use some data structure, such as a list, to store the contents of M 's tape
- Now you can simulate M step by step, updating the data structure as you move along
- The **interpreter** will continue the simulation until the machine eventually halts

The Existence of a Universal Turing Machine

- How would you code U in your favorite programming language?
- First, you need to decide on some representation scheme for M .
 - For example, you can use an array or a dictionary to encode M 's transition function
- Then you would use some data structure, such as a list, to store the contents of M 's tape
- Now you can simulate M step by step, updating the data structure as you move along
- The **interpreter** will continue the simulation until the machine eventually halts
- Translating the interpreter above into the corresponding Turing machine is “easy”

Universal Turing Machine: Implications

- There is more than one Turing machine U that works as indicated above

Universal Turing Machine: Implications

- There is more than one Turing machine U that works as indicated above
- The existence of even a *single* such machine is already fundamental to computer science

Universal Turing Machine: Implications

- There is more than one Turing machine U that works as indicated above
- The existence of even a *single* such machine is already fundamental to computer science
- The idea of a “universal program” is of course not limited to theory

Universal Turing Machine: Implications

- There is more than one Turing machine U that works as indicated above
- The existence of even a *single* such machine is already fundamental to computer science
- The idea of a “universal program” is of course not limited to theory
- The most famous practical example is represented by **compilers** (for programming languages), which are often used to compile themselves!

Summary

- We have discussed Turing machines (TMs) as the standard **model of computation**

Summary

- We have discussed Turing machines (TMs) as the standard **model of computation**
- TMs and every other computational model independently proposed have all the same power (**Church-Turing thesis**)

Summary

- We have discussed Turing machines (TMs) as the standard **model of computation**
- TMs and every other computational model independently proposed have all the same power (**Church-Turing thesis**)
- Computable functions (total/partial) are those which can be computed by a TM

Summary

- We have discussed Turing machines (TMs) as the standard **model of computation**
- TMs and every other computational model independently proposed have all the same power (**Church-Turing thesis**)
- Computable functions (total/partial) are those which can be computed by a TM
- There exists few variants of standard TM like multi-tape or non-deterministic TMs yet they all have the same power

Summary

- We have discussed Turing machines (TMs) as the standard **model of computation**
- TMs and every other computational model independently proposed have all the same power (**Church-Turing thesis**)
- Computable functions (total/partial) are those which can be computed by a TM
- There exists few variants of standard TM like multi-tape or non-deterministic TMs yet they all have the same power
- The existence of a special Universal Turing Machine (UTM) allows us to design an algorithm that can run any other algorithm